

# Redis 源码剖析



# 前言

Redis(Remote Dictionary Server)

- Key-Value NoSQL
- 内存 亚毫秒级访问
- 简单的API

源码面前，了无秘密 -- 《STL源码剖析》侯捷

本次分享从源码入手去分析Redis后端的精巧设计

源码为 Redis v3.2.12 (Tencent Dev-Cloud yum 源版本)

涉及到 Redis v6.2.4 (Redis Stable 最新版本)

# 目录

- Redis I/O : Reactor 模式
- Redis DB
- Redis 持久化 : AOF / RDB

## Reactor 模式 -- 起源

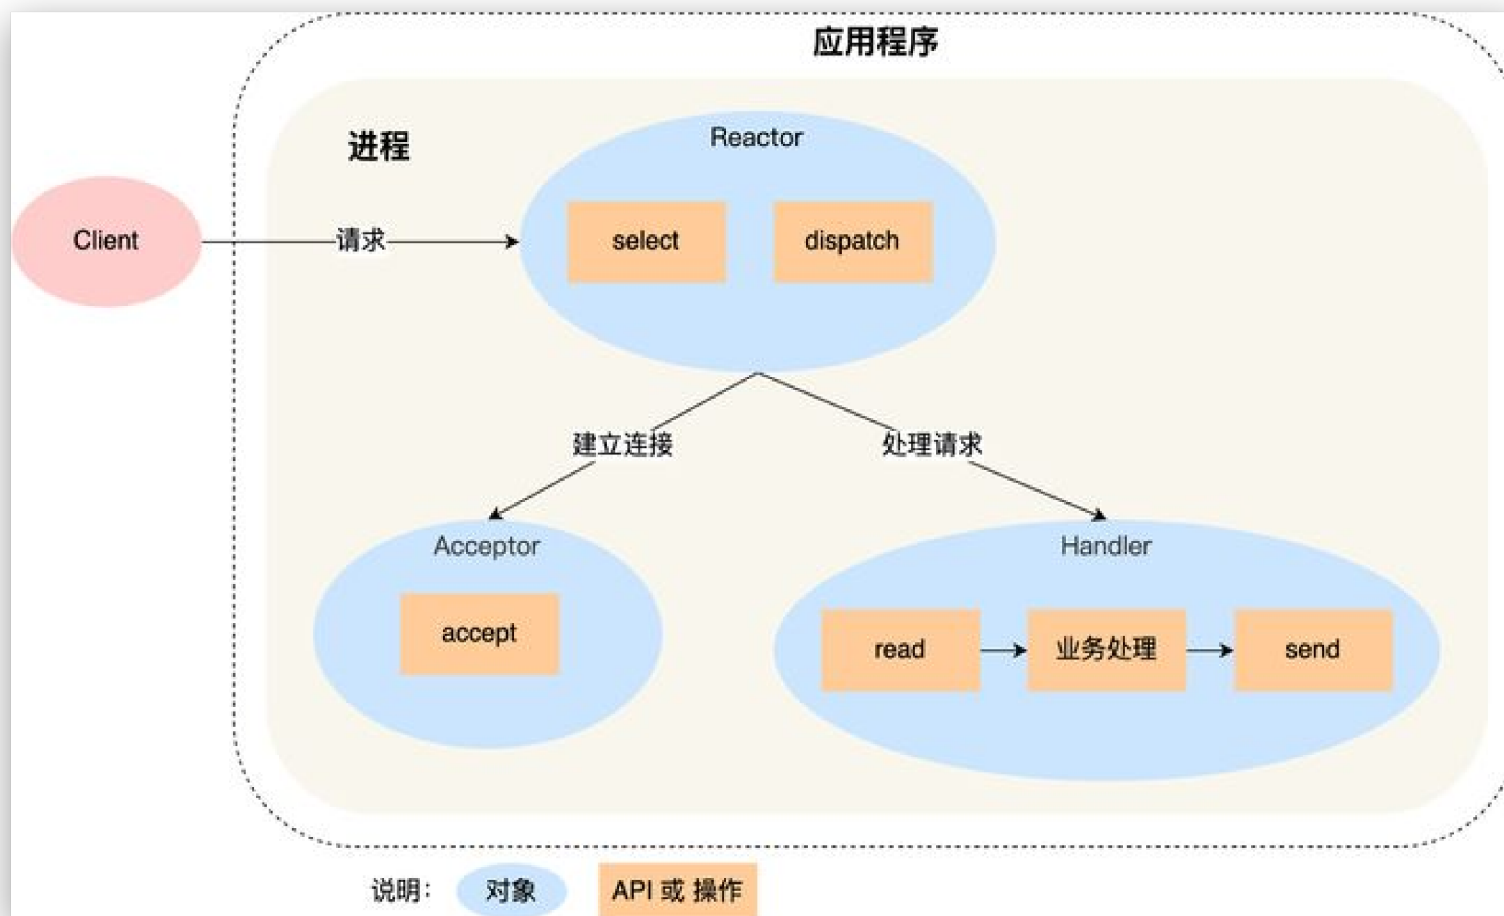
Client / Server结构

redis/3.2.12/src/server.h

```
struct redisServer {  
    /* General */  
    // ...  
    redisDb *db;  
    list *clients;                /* List of active clients */  
  
    /* Networking */  
    // ...  
    int ipfd[CONFIG_BINDADDR_MAX]; /* TCP socket file descriptors */  
    int ipfd_count;                /* Used slots in ipfd[] */  
}
```

多进程？ 多线程？ 线程池？ I/O多路复用？ 面向过程编程？ 面向对象编程？

## Reactor 模式 -- 架构



## Reactor 模式 -- 实现(1)

redis/3.2.12/src/ae.h

```
#ifdef HAVE_EVPORT
#include "ae_evport.c"
#else
    #ifdef HAVE_EPOLL
        #include "ae_epoll.c"
    #else
        #ifdef HAVE_KQUEUE
            #include "ae_kqueue.c"
        #else
            #include "ae_select.c"
        #endif
    #endif
#endif
#endif
```

## Reactor 模式 -- 实现(2)

redis/3.2.12/src/ae.c

```
aeEventLoop *aeCreateEventLoop(int setsize) {}

int aeCreateFileEvent(aeEventLoop *eventLoop, int fd, int mask,
                      aeFileProc *proc, void *clientData) {}

void aeDeleteFileEvent(aeEventLoop *eventLoop, int fd, int mask) {}

int aeProcessEvents(aeEventLoop *eventLoop, int flags) {}

void aeMain(aeEventLoop *eventLoop) {}

void aeSetBeforeSleepProc(aeEventLoop *eventLoop, aeBeforeSleepProc *beforesleep) {}
```

## Reactor 模式 -- 单Reactor瓶颈

跑满 CPU 一个核心，达到瓶颈  
单机 QPS 10W



## Reactor 模式 -- Redis 6 Multi-Reactor

redis/6.2.4/src/networking.c

```
void initThreadedIO(void) {  
    // ...  
    /* Spawn and initialize the I/O threads. */  
    for (int i = 0; i < server.io_threads_num; i++) {  
        // ...  
        pthread_mutex_init(&io_threads_mutex[i], NULL);  
        pthread_mutex_lock(&io_threads_mutex[i]); /* Thread will be stopped. */  
        if (pthread_create(&tid, NULL, IOThreadMain, (void*)(long)i) != 0) {  
            serverLog(LL_WARNING, "Fatal: Can't initialize IO thread.");  
            exit(1);  
        }  
        io_threads[i] = tid;  
    }  
}
```

## Redis DB -- 键空间

redis/3.2.12/src/server.h

```
typedef struct redisDb {  
    dict *dict;                /* The keyspace for this DB */  
    dict *expires;             /* Timeout of keys with a timeout set */  
    int id;                    /* Database ID */  
    // ...  
} redisDb;  
  
typedef struct dictEntry {  
    void *key;  
    union {  
        void *val;  
        // ...  
    } v;  
    struct dictEntry *next;  
} dictEntry;
```

## Redis DB -- 过期机制

惰性 (CPU友好) redis/3.2.12/src/server.c

```
int expireIfNeeded(redisDb *db, robj *key) {
    mstime_t when = getExpire(db, key);
    mstime_t now;
    if (when < 0) return 0; /* No expire for this key */
    // ...
    return dbDelete(db, key);
}
```

定期 (内存友好) redis/3.2.12/src/server.c

```
int activeExpireCycleTryExpire(redisDb *db, dictEntry *de, long long now) {
    long long t = dictGetSignedIntegerVal(de);
    if (now > t) {
        // ...
    }
}
```

## Redis DB -- 对象与编码

```
typedef struct redisObject {
    unsigned type:4;
    unsigned encoding:4;
    void *ptr;
    // ...
} robj;

#define OBJ_STRING 0
#define OBJ_LIST 1
#define OBJ_SET 2
#define OBJ_ZSET 3
#define OBJ_HASH 4

#define OBJ_ENCODING_RAW 0      /* Raw representation */
#define OBJ_ENCODING_INT 1      /* Encoded as integer */
#define OBJ_ENCODING_HT 2      /* Encoded as hash table */
#define OBJ_ENCODING_ZIPMAP 3   /* Encoded as zipmap */
#define OBJ_ENCODING_LINKEDLIST 4 /* Encoded as regular linked list */
#define OBJ_ENCODING_ZIPLIST 5   /* Encoded as ziplist */
#define OBJ_ENCODING_INTSET 6    /* Encoded as intset */
#define OBJ_ENCODING_SKIPLIST 7  /* Encoded as skiplist */
#define OBJ_ENCODING_EMBSTR 8    /* Embedded sds string encoding */
#define OBJ_ENCODING_QUICKLIST 9 /* Encoded as linked list of ziplists */
```

## Redis DB -- 猜猜编码是啥？

```
// object : string
SET msg1 "hi"
SET msg2 123

// object : hash
HSET book name "Mastering Cpp in 21 days"

// object : zset
ZADD price 10.0 t-shirt 5.0 shoes
```

## Redis DB -- ziplist

redis/3.2.12/src/ziplist.h

```
typedef struct zentry {  
    unsigned int prevrawlen, prevrawlen;  
    unsigned int len, len;  
    unsigned int headersize;  
    unsigned char encoding;  
    unsigned char *p;  
} zentry;
```

redis/3.2.12/src/dict.h

```
typedef struct dict {  
    dictType *type;  
    dictht ht[2];  
    long rehashidx; /* rehashing not in progress if rehashidx == -1 */  
    // ...  
} dict;
```

## Redis 持久化 : RDB (Redis DataBase)

redis/3.2.12/src/rdb.c

```
int rdbSave(char *filename) {}  
int rdbLoad(char *filename) { \\ ... }
```

## Redis 持久化 : AOF (Append Only File)

```
struct redisServer {  
    // ...  
    sds aof_buf;  
}
```

```
# appendfsync always  
appendfsync everysec  
# appendfsync no
```



# Redis 持久化: AOF重写

redis/3.2.12/src/aof.c

```
int rewriteAppendOnlyFileBackground(void) {
    if ((childpid = redisFork(CHILD_TYPE_AOF)) == 0) {
        /* Child */
        // ...
        if (rewriteAppendOnlyFile(tmpfile) == C_OK) {
            sendChildCowInfo(CHILD_INFO_TYPE_AOF_COW_SIZE, "AOF rewrite");
            exitFromChild(0);
        } else {
            // ...
        }
    } else {
        /* Parent */
        // ...
        replicationScriptCacheFlush();
    }
}
```